

Why I Believe the REGAL Trial Will Fail and SLS Will Crash

Notes on galinpepimut-S, SELLAS Life Sciences (SLS), and the Phase 3 REGAL trial in AML

by Franciszek Kessler

July 2026 — personal research notes, not investment advice

Table of Contents

The base rate for this entire drug class is close to zero	3
SELLAS has already run this exact experiment once, and it failed	3
The biology is harder than the target rationale suggests	4
The Phase 2 data isn't the slam dunk it's presented as	5
The trial itself has some structural soft spots	6
Why a delayed trial is not the bullish signal it looks like	7
<i>The more likely explanation: the original median OS assumption was wrong</i>	7
<i>The healthy-volunteer / trial-selection effect</i>	7
<i>Documented, named examples of exactly this pattern</i>	7
<i>The pattern shows up repeatedly in small-cap biotech, too</i>	8
The immune response data doesn't hold up as well as it's presented	9
Valuation: the market is already pricing this as a win	9
<i>Sizing the actual commercial opportunity</i>	9
<i>What does that mean for the current price</i>	1
<i>Dilution, insider incentives, and thin coverage</i>	11
<i>A crowded short, too</i>	11
<i>Knowing who the fool at the poker table is</i>	11
Where I land	12
<i>Where the stock actually goes if I'm right</i>	12
Sources	14

I don't think REGAL is going to read out the way the bulls expect, and the stock's run-up into the data has gotten ahead of the evidence. Below is the case, roughly in the order I'd walk someone through it, with the sources at the end. I've also built a companion Excel model with my full valuation and TAM build-out; the numbers referenced in the valuation section below come from that model.

The base rate for this entire drug class is close to zero.

Start with the modality itself, before even getting to GPS specifically. Peptide vaccines that target tumor-associated antigens (TAAs) — proteins that are overexpressed in cancer but not mutated, like WT1 — have never produced a single FDA approval. Not one. An industry analysis tracking synthetic peptide vaccine candidates found that of more than 1,000 examined, 125 reached Phase I and 30 reached Phase II, and none has ever cleared Phase III to reach the market.

People sometimes point to sipuleucel-T (Provenge) as a counterexample, but that's actually a cell-based therapy — dendritic cells extracted from the patient and pulsed with a fusion protein outside the body — not a standalone injected peptide vaccine like GPS. Mechanistically, it's a different class of drug with a different manufacturing and regulatory path, so it doesn't really count against the pattern.

The literature is consistent on why this keeps happening: an immunosuppressive tumor microenvironment, tumor immune evasion, and tumor heterogeneity that lets low-antigen subclones survive selection. None of that is specific to a bad target choice — it's a property of the approach, and it doesn't go away because WT1 happens to be a well-validated target on paper.

Two prior examples make this concrete. Rindopepimut, a peptide vaccine for glioblastoma, had three positive Phase 2 trials, an FDA breakthrough designation, and a strong measured immune response — and then its placebo-controlled Phase 3 (ACT IV) was stopped early for futility. Tecemotide (Stimuvax), the MUC1 vaccine you and I already looked at, missed its Phase 3 primary endpoint in lung cancer despite years of buildup, and its own Phase 2 immune response rate was only around 21%. Both are close structural analogs to GPS: same broad mechanism, same encouraging early data, same outcome.

SELLAS has already run this exact experiment once, and it failed

This is the part that gets underweighted. SELLAS's other flagship drug, nelipepimut-S (NeuVax), is a HER2-targeting peptide vaccine using the same platform — a short peptide plus GM-CSF adjuvant — but aimed at breast cancer rather than AML. Its pivotal Phase 3 (the PRESENT trial, 758 patients, randomized, double-masked, placebo-controlled) was stopped early for futility. The independent monitoring committee reviewed the interim data and concluded that there was no realistic path to a positive result.

SELLAS kept the program alive afterward in narrower settings — a DCIS trial (VADIS), a terminated Phase 2b combination study, and some FDA discussions about a triple-negative breast cancer indication that didn't progress to an actual registration trial. So the pattern with this company's other big asset is: strong early data, a clean pivotal failure, then a retreat into smaller subpopulations rather than a full stop.

On management specifically, Angelos Stergiou has run SELLAS since founding it in 2012. Before that, his C-suite and senior operational roles were at Genesis Life Sciences, PAION, Accentia Biopharmaceuticals, BioVest International, Analytica International, and Anavex Life Sciences, roughly across 2002–2009. I couldn't find a drug he led to FDA approval at any of these companies or at SELLAS — there's no track record of a win to weigh against NeuVax's failure.

The Anavex connection is worth sitting with for a moment, separate from the NeuVax point. Anavex's flagship drug, blarcamesine (ANAVEX2-73), has developed a genuinely rough public record in the years since its Rett syndrome Phase 2/3 trial (EXCELLENCE) missed its endpoints in January 2024. Its Alzheimer's Phase IIb/III trial has drawn a formal published complaint from outside researchers alleging inconsistent reporting of results over time, an undisclosed protocol and statistical analysis plan, and a co-primary endpoint (ADCS-ADL) that didn't reach significance even though the company's public messaging emphasized the endpoint that did. One independent report went further, describing a pattern of “spinning weak or failed clinical trial data into headline-friendly fluff.” I want to be careful here: Stergiou's time at Anavex was early, well before any of this unfolded, and I have no evidence he was personally involved in the conduct being criticized. But it's still a second company on his résumé, alongside SELLAS/NeuVax, where the flagship asset's public data story and the underlying trial results don't fully line up — and that's the kind of pattern worth just noticing rather than explaining away.

The biology is harder than the target rationale suggests

WT1 looks great on paper — it's overexpressed in nearly all AML cases and is mostly absent from normal adult tissue, which is why an NCI panel ranked it the top-priority cancer antigen back in 2009. But a good target on paper and a good vaccine target are different things, and WT1 runs into the same wall most self-antigens do.

WT1 is a normal protein your body already makes, not a mutated, foreign-looking one. The immune system spends your whole life deleting or muting T cells that react strongly to self-proteins, specifically to prevent autoimmunity — so a lot of the T cells you'd need for a powerful anti-WT1 response were never available to begin with. This is the textbook reason TAA vaccines underperform, and it's why the field has been moving toward neoantigen vaccines (mutation-derived, genuinely foreign) instead. On top of that, a [2004 Journal of Clinical Investigation study](#) found a second, independent problem: some self/tumor peptides don't sit tightly enough in the MHC groove on the cell surface to hold a T cell's attention long enough to activate it — a stability issue, separate from tolerance.

There's also a straightforward escape route. AML is genetically messy — lots of subclones, not one uniform population. Put a WT1-vaccine-trained immune system under selection pressure, and the leukemic cells that already express less WT1, or that downregulate the MHC machinery needed to display it, have a survival advantage and can repopulate the disease without ever needing a new mutation. One study specifically looking at post-vaccination relapses didn't find WT1 loss or MHC downregulation driving them, which is mildly reassuring. Still, it's a small dataset with short follow-up, so I wouldn't lean on it too hard.

It's also worth noting that WT1 sits inside the cell (it's a transcription factor), so ordinary antibodies can't reach it at all — that's exactly why this has to be a T-cell vaccine approach rather than an antibody

drug. People have tried building TCR-mimic antibodies against the WT1-peptide-MHC complex instead. Still, none of that work has made it past the preclinical stage, which suggests this target has resisted more than one attempted solution, not just this one.

The Phase 2 data isn't the slam dunk it's presented as

Most of the bull case ultimately traces back to GPS's Phase 2 track record, so it's worth putting that data under the same scrutiny I'd apply to anything else in this thesis. It's weaker support than it's usually presented as, for three separate reasons.

First, the foundational trial (Maslak et al., Blood Advances, 2018; NCT01266083) was a single-arm, non-randomized trial. There was no concurrent control group — patients on GPS were compared against a historical benchmark drawn from older literature (a 20–25% expected 3-year OS rate), not against a contemporaneous group of similar patients followed at the same time, at the same centers, under the same standard of care. That's exactly the historical-control-bias problem I've flagged elsewhere in this document as a weakness of the REGAL comparator — except here it sits inside the primary piece of evidence used to justify running REGAL in the first place. A single-arm trial beating an historical control is real evidence of something. Still, it's a much weaker design than a randomized comparison, and it's the kind of design that gets a drug into a randomized Phase 3 rather than the kind that reliably predicts what that Phase 3 will show.

Second, and I think this is the more important point: that trial enrolled patients in CR1 — first complete remission — not CR2, which is the population REGAL is actually testing. CR1 patients haven't relapsed yet; they're a meaningfully healthier, better-prognosis group than CR2 patients who have already relapsed once and gone through salvage therapy. Encouraging survival data in a first-remission population doesn't necessarily transfer to a second-remission population with a different, worse baseline prognosis. Citing the CR1 data as support for the CR2 pivotal trial is, in a meaningful biological sense, comparing results across two different diseases, even though both are labeled "AML in remission."

Third, the trial was tiny — 22 patients enrolled, and the widely cited 47% 3-year survival figure is based on just 19 evaluable patients. At that size, each patient's outcome is worth roughly 5 percentage points of the headline statistic. A single additional long-term survivor, or a single additional early relapse, would meaningfully move the number in either direction. The "median OS not reached, expected to exceed 67.6 months" framing sounds dramatic. Still, with a cohort this small, a median not yet reached is also just what you'd expect from too few events in too few patients — it doesn't have to reflect a large true effect size. The confidence interval around that number is almost certainly very wide, even though it isn't usually reported alongside the headline figure.

There's also a separate, randomized CR2 Phase 2 comparison that's publicly cited (21.0 vs. 5.4 months OS) and that I've referred to earlier in this document as the strongest piece of supporting evidence, since it's randomized rather than single-arm. I want to flag some uncertainty about that data point directly: I have not been able to locate a clear, independently verifiable, peer-reviewed publication describing that specific randomized comparison with the same level of methodological detail as the CR1 paper above. It may exist and be less prominently indexed. Still, as things stand, the single most-cited

piece of "randomized" support for GPS is, from what I can verify, less rigorously documented than the single-arm CR1 data, which is a strange asymmetry, since the single-arm data is the weaker design but the better-sourced one. I'd hold the randomized CR2 Phase 2 claim more loosely than it's typically presented, and I'd treat "a randomized supporting Phase 2 with a large effect size" — one of the genuine positives I list at the end of this document — as somewhat less certain than the way I originally framed it.

The trial itself has some structural soft spots.

REGAL is randomizing 127 transplant-ineligible CR2 patients 1:1 to GPS or the investigator's choice of best available therapy (BAT), open-label, with a final analysis triggered at 80 pooled deaths. A few things about that design give me pause.

- 127 patients and 80 events is a small trial for a definitive Phase 3 — there's real room for a modest true effect to land on either side of significance just from chance.
- It's open-label, not blinded. That tends to introduce soft biases — such as follow-up frequency and the threshold for starting rescue therapy — that generally inflate the treated arm's apparent benefit.
- The trial's historical benchmark (roughly 6 months expected survival for this population) may be stale. Venetoclax-based regimens and oral azacitidine both became standard of care in 2020, right as REGAL launched in February 2021 and enrolled through 2024. If the modern BAT arm is doing better than that old assumption, it dilutes GPS's apparent benefit even if GPS is doing exactly what it did in Phase 2. Against that, venetoclax-based therapy in the closest comparable setting (active relapsed/refractory AML) still shows only around 6.1 months median survival, so this risk is real but not conclusively proven.
- There's a statistical multiplicity risk worth watching for at readout. REGAL has secondary endpoints beyond overall survival — leukemia-free survival, time to AML progression. If the primary OS endpoint narrowly misses, watch for the press release framing to lean on whichever secondary endpoint looks best; that's a well-worn move when a primary endpoint disappoints, and it doesn't change what the trial was actually designed and powered to show.
- Even in a scenario where the trial succeeds, there's residual manufacturing risk between readout and launch. Peptide-plus-Montanide-adjuvant vaccines have historically had real-world batch consistency and supply issues in other programs; that's a separate, unresolved risk sitting between a positive trial result and a smooth commercial launch, and it doesn't show up in a topline hazard ratio.

I've broken the timing question — whether REGAL taking longer than planned is good news or bad news — out into its own section below, because I think it deserves a full treatment rather than a bullet point.

Why a delayed trial is not the bullish signal it looks like

A lot of the bull case for GPS leans on the fact that REGAL has taken longer than originally planned to reach its event threshold, on the theory that if patients are living longer, the drug must be working. This reasoning has the mechanism backward more often than not, and it's worth spelling out why in detail rather than just asserting it.

The more likely explanation: the original median OS assumption was wrong

REGAL is an event-driven trial — it needs 80 pooled deaths (across both arms combined) to trigger final analysis, regardless of how long that takes in calendar time. As of the writing of this paper we are at 78 out of 80 events. When a trial like this takes longer than planned, the arithmetic reason is always the same: fewer people are dying per month than the original design assumed. The question is why. There are really only two possibilities, and they point in opposite directions:

- GPS patients specifically are living much longer than assumed, which is the bullish story; or
- The pooled population — which includes the control arm — is living longer than assumed for reasons that have nothing to do with GPS: a stale historical benchmark. This control arm benefits from newer supportive-care drugs that didn't exist when the benchmark was set, or from a trial population that is inherently healthier than the historical patients from whom the benchmark was built.

That second explanation — a miscalculated starting assumption, not a working drug — has much stronger empirical support behind it as the typical cause. A systematic review of Phase III oncology trials from 2013 to 2023, published in ESMO Open, focused on trials that reached full enrollment but encountered statistical power issues. It found that a lower-than-expected event rate caused 71.2% of those power failures, and the dominant reason across the trials studied was that the control arm survived longer than the assumption used to power the study, not that the experimental drug outperformed expectations.

The healthy-volunteer / trial-selection effect

There's a second, compounding reason trial populations tend to outlive historical benchmarks that has nothing to do with drug efficacy at all: patients who enroll in clinical trials are not a random sample of all patients with a given disease. Trials have eligibility criteria — adequate organ function, a minimum performance status, and no major uncontrolled comorbidities — that systematically screen out the sickest, frailest patients who would otherwise pull down a historical survival benchmark. A trial population, on both arms, will tend to outlive an unselected historical cohort simply because of who was allowed to enroll, independent of any treatment effect. Combine that with a genuine improvement in supportive care since the benchmark was set, and a pooled survival curve can end up looking dramatically better than the historical assumption without either arm's actual treatment having anything to do with it.

Documented, named examples of exactly this pattern

This isn't a hypothetical concern — there are concrete, publicly documented trials where exactly this happened, and the sponsors themselves had to file protocol amendments acknowledging it:

- A Phase 3 adjuvant everolimus trial in kidney cancer (S0931) had to amend its protocol after actual median recurrence-free survival in the control arm came in at 6.6 years versus the 6-year assumption used at design — the control arm did better than planned, slowing event accrual and forcing the trial to add years of additional follow-up.
- A head-and-neck cancer trial testing radiation plus cetuximab after surgery saw a similar problem: its own protocol amendment attributes a multi-year delay directly to a lower-than-expected event rate, pushing the final analysis back by years.
- A pancreatic cancer adjuvant chemotherapy trial hit the same wall — a lower-than-expected event rate pushed the projected final analysis out to 2029 or later, and the trial's statisticians had to formally accept reduced statistical power rather than wait indefinitely for events to accumulate.

A secretly brilliant experimental drug was driving none of these three trials. In each case, the documented cause was a miscalculated baseline assumption about how the control arm — and the trial population generally — would behave, not a treatment effect. This is the same mechanism, I think, that is most likely operating in REGAL: supportive care for AML (venetoclax combinations, oral azacitidine) improved substantially right around when REGAL launched and enrolled, and REGAL's population, like any trial population, was screened for eligibility in ways that make it healthier than the historical cohort its ~6-month benchmark was drawn from.

The pattern shows up repeatedly in small-cap biotech, too.

Beyond the systematic academic evidence, there's a well-worn pattern: small-cap oncology stocks run up specifically when a trial is taking a long time, on the theory that the delay itself is good news — and it has an almost uniformly bad track record. Telik's canfosfamide failed two separate Phase 3 ovarian cancer trials after years of anticipation, with one trial showing literally identical median progression-free survival (3.5 months) in both the drug arm and the control arm. GenVec's TNFerade is a particularly sharp illustration of the risk in reading an early positive-looking number too optimistically: a first interim analysis suggested roughly a 25% lower risk of death, but by the next scheduled look, months later, that had shrunk to a statistically meaningless 8% (hazard ratio 0.921), and the company scrapped the trial entirely. Novelos' NOV-002 missed its Phase 3 endpoints in lung cancer despite encouraging earlier data. Cell Therapeutics' pixantrone showed a real effect, but it wasn't enough to satisfy an FDA advisory panel, which voted unanimously against approval.

Compare all of that to what actually happens when a drug's effect is genuinely large: Roche's vemurafenib (Zelboraf) in melanoma, Onyx's sorafenib (Nexavar) in kidney cancer, Algeta's radium-223 (Alpharadin) in prostate cancer, and Johnson & Johnson's abiraterone (Zytiga) in prostate cancer were all stopped early, specifically because their independent monitoring committees saw a survival benefit clear and large enough that continuing to randomize patients to the inferior control arm became both statistically unnecessary and ethically indefensible. Big, real drug effects tend to get caught quickly. Trials that quietly run long, past their planned timeline, are empirically much more often explained by a miscalculated starting assumption than by a hidden winner.

To be fair to the other side of this specific point: REGAL's January 2025 interim analysis did involve an independent committee reviewing unblinded data and confirming GPS exceeded the pre-specified

futility threshold. That is a real, if modest, piece of positive information that none of the failure examples above had at a comparable stage. It doesn't override the base rate, but it means the delay isn't purely a bad sign here — it's just not the confirmation of success I think a lot of the market is treating it as.

The immune response data doesn't hold up as well as it's presented.

The same skepticism that applies to Stimuvax's 21% immune response rate applies here. In the foundational GPS Phase 2 trial in CR1 AML ([Maslak et al., Blood Advances, 2018](#), n=22), only 14 of the 22 enrolled patients had any immune response data collected at all — 8 were never tested, some likely because they relapsed and dropped out early. Of the 14 tested, 9 were positive in at least one assay, which is where the commonly cited “64%” figure comes from. But 9 out of the full 22-patient population is 41%, not 64%, and the gap between those two numbers is exactly the kind of thing worth noticing.

The CD8+ killer T-cell response — arguably the more important one mechanistically — could be measured only in HLA-A*02:01-positive patients, which accounted for 9 of the 22 enrolled. Most of the trial population couldn't even be evaluated on that measure. And when the authors checked whether immune responders actually did better than non-responders on survival, the difference didn't reach statistical significance ($p=.11$ and $p=.08$), and one patient with a confirmed T-cell response still relapsed anyway.

The more recent, better-sounding numbers — 64% in the randomized CR2 Phase 2, 80% in the REGAL interim — come from press releases and conference disclosures rather than a published paper with a disclosed sample size or per-assay breakdown. The trend has been toward more impressive headline numbers with less available methodology to verify them, which is the opposite of what I'd want to see.

Valuation: the market is already pricing this as a win

This is the part that ultimately matters most for the stock, independent of everything above. As of July 5, 2026, SLS is trading at \$14.98/share. With roughly 186 million shares outstanding, the market cap is about \$2.8 billion. Against roughly \$107 million of cash and no debt, that's an enterprise value of about \$2.7 billion for a company whose lead asset does not yet have a positive Phase 3 readout, let alone an FDA approval.

Sizing the actual commercial opportunity

Working through the addressable patient population for GPS in its lead indication: of roughly 20,000 new AML cases diagnosed annually in the U.S., perhaps 12,000 reach a first complete remission, of whom a majority eventually relapse, and only a minority of those who relapse and receive salvage therapy go on to reach a second complete remission — one dataset put that CR2 rate at roughly 36% of relapsed patients. Removing the share who are transplant-eligible (REGAL's population is transplant-ineligible only) leaves a rough estimate of 1,000–1,300 new eligible U.S. patients per year, or perhaps 1,500–2,000 at steady state, once you account for patients remaining on therapy for a period of time. Extending that to other developed markets outside Greater China (SELLAS licensed China rights separately) yields a rough global addressable population of 3,000–4,500 patients at steady state.

At a plausible orphan-drug price in line with other targeted AML therapies (roughly \$150,000–\$200,000 per year, comparable to venetoclax or gilteritinib), that implies a peak, fully-penetrated total addressable market of roughly \$450 million to \$900 million annually — and realistic peak sales, assuming a more typical 30–50% share of that addressable population rather than full penetration, in the rough range of \$150 million to \$450 million per year.

A few factors could push that realistic range even lower. GPS is licensed from Memorial Sloan Kettering, and I haven't been able to confirm exactly how many years of composition-of-matter protection remain once you layer that on top of standard orphan drug exclusivity (7 years from approval) — the actual monetizable window could be shorter than a simple peak-sales model assumes. Reimbursement is a separate friction: payers scrutinize a \$150,000–\$300,000-a-year therapy most harshly exactly when the underlying survival benefit is modest, and the population is narrow, which is GPS's situation if it clears the trial at all. And the competitive landscape isn't static — menin inhibitors like revumenib and other emerging AML maintenance approaches could carve into the same eligible population by the time GPS could plausibly launch, further capping the realistic addressable market from the demand side, not just the supply side.

What does that mean for the current price?

An enterprise value of \$2.7 billion against a realistic peak sales estimate of \$150–450 million implies a multiple of roughly 6–18x realistic peak sales — for a single indication, pre-approval, ahead of a trial result I'd estimate has something like a 5% chance of even hitting statistical significance, before getting to the separate question of whether a narrow win translates into a commercially compelling, broadly reimbursed drug. Multiples like that are sometimes justified for a drug that has already been de-risked by an approval and is showing an early launch trajectory. They are very hard to justify for an asset still awaiting its pivotal readout, for a drug class with a zero-for-history approval record, from a company whose only other pivotal trial on this exact platform failed. My read is that the current price isn't really reflecting a 5% probability-weighted outcome — it's closer to what you'd expect if the market had already assumed the trial succeeds, and the drug launches roughly on plan. That's the asymmetry I'd want anyone looking at this name to sit with: the downside case (trial disappoints, EV re-rates toward cash-plus-optionality) is a much larger percentage move than the upside case (trial succeeds, and the stock has to re-rate from a starting point that already assumed a lot of that success).

Worth stress-testing this with a genuinely generous bull case, too, because even the best-case math doesn't get you very far from here. Assume 4,500 patients (the high end of the global addressable population), a \$300,000 annual cost per patient (well above the \$150,000–\$200,000 comparable range I used above), a 3% annual price increase to keep pace with inflation, and extremely fast adoption sustained over 15 years. That gets you to roughly \$2 billion to \$2.5 billion in peak annual sales — genuinely blockbuster territory. Layer on a generous 40% profit margin and discount the resulting cash flows back, and you land on an NPV of roughly \$5.4 billion, or about \$29 a share. That's a double from today's \$14.98 — which sounds attractive until you notice what had to go right to get there: extremely fast, near-total market adoption, a premium price nearly double what comparable AML drugs command, sustained pricing power for 15 straight years, and a 40% margin on top of all of it. That's a miracle commercial scenario, stacked on top of the separate miracle of the trial actually working in the

first place (itself, by my estimate, only about a 5% probability event). Multiplying a generous bull-case commercial outcome by a roughly 1-in-20 probability of the underlying trial succeeding is not a combination that supports paying today's price for a double as the best realistic case.

Dilution, insider incentives, and thin coverage

A few structural points compound the valuation concern. Shares outstanding have grown 86.4% over the past year — dilution on that scale, largely through warrant exercises, means any per-share value in this document is being measured against a moving denominator. There's more room to come: SELLAS established an at-the-market (ATM) facility of up to \$150 million in 2026 that hadn't been used as of the most recent quarterly update, which is a real overhang sitting on top of the current share count, regardless of trial outcome. Separately, CEO Angelos Stergiou received a grant of 1,065,000 restricted stock units on June 16, 2026, just weeks ahead of the anticipated readout, lifting his direct holdings to 1,759,183 shares. That's a compensation award, not an open-market sale, so it isn't the classic insider-dumping red flag — but a grant of that size timed immediately before a binary catalyst is still worth noting as a governance and incentive-alignment data point. Finally, the stock is covered by only 2–4 analysts, which is thin even for a small-cap biotech — it means there's relatively little independent institutional scrutiny of the name. It helps explain why retail sentiment (see below) can end up dominating the public narrative around the stock rather than being counterbalanced by sell-side coverage.

One fairness point on the other side: SELLAS also has a second pipeline asset, SLS009 (tambiciclib), a CDK9 inhibitor with positive Phase 2 signals in relapsed/refractory AML presented at ASH 2025 and now dosing in a newly diagnosed frontline AML study. Some portion of the current enterprise value plausibly reflects optionality on SLS009, not a pure bet on GPS/REGAL — I've priced the EV against GPS-only peak sales above, which somewhat overstates how purely GPS-dependent the stock is. That said, SLS009 is still early (Phase 2, unpartnered, in a crowded CDK9 field), and I don't think it comes close to closing the gap between today's price and a probability-weighted view of the REGAL outcome on its own.

A crowded short, too

One more data point worth flagging: the stock borrow rate on SLS has reportedly been running around 150% — an unusually high cost to borrow shares to short. That tells you demand to short the name is intense relative to the available float, which cuts two ways. It confirms there's a well-populated bear case already being expressed in the market (I'm clearly not alone in this view), but a hard-to-borrow stock with a lot of short interest is also exactly the setup that can produce a sharp, mechanical short squeeze on any positive news, regardless of the news' actual durability or quality. Worth factoring into position sizing and timing, separate from the fundamental thesis itself.

Knowing who the fool at the poker table is

There's an old poker line: if you've been at the table for thirty minutes and don't know who the fool is, it's you. I'd apply that here. There are two dedicated retail communities built entirely around this stock, [r/SLS Stock](#) and [r/sellaslifesciences](#), and I think the mere existence of a fanbase around a single small-cap biotech ahead of a binary trial readout is itself a warning sign, independent of anything discussed in

either forum. Reading through the posts, the focus is overwhelmingly on price action and bullish framing rather than on the underlying clinical or statistical facts — chart patterns, short-interest gossip, and speculation about what price the stock will hit — rather than on serious engagement with trial design, historical base rates, or the mechanism. One post that stuck with me: a poster said that if SELLAS reaches \$40 billion, he'll name his son GPS. That's not a data point about the drug. It's a data point about the crowd holding the stock and how much of the current price is driven by narrative and community enthusiasm rather than a probability-weighted assessment of the actual evidence. A stock with an organized retail fanbase, a 150% borrow rate, and a valuation that already assumes success is a stock where it's worth asking, plainly, who's on the other side of your trade — and whether it's you.

Where I land

Weighing the class-wide zero-approval history, SELLAS' own prior Phase 3 failure on the same platform, the tolerance and antigen-escape problems specific to WT1, weaker-than-presented Phase 2 support (single-arm, wrong remission stage, small sample, and a randomized CR2 comparison I can't fully verify), a small open-label trial with a possibly stale control-arm assumption, the empirical pattern linking slow event accrual to underpowered or negative trials, immune-response data that's weaker than its headline framing suggests, and a valuation that looks like it already assumes success — against the genuine positives that remain, namely a weak historical comparator and an interim analysis that did clear its futility bar — I land at around a 5% probability of REGAL hitting its primary endpoint. Not a coin flip and not even close to one — a long shot. The genuine positives keep this from being a flat zero. Still, they don't come close to overriding a drug class with no approvals ever, a same-company, same-platform prior failure, and a mechanism that fights self-tolerance on top of everything else. This is dramatically below what I think the current share price assumes.

Where the stock actually goes if I'm right

If REGAL fails, GPS is effectively worthless as a going concern, and the stock should trade down toward whatever is left: cash, plus whatever the market is willing to pay for the remaining pipeline. Cash and equivalents were \$107.1 million as of March 31, 2026. Against roughly 186 million shares outstanding, that's about \$0.58 per share in cash alone — and that \$0.58 is itself a ceiling, not a floor, since a failed pivotal trial would likely trigger cost cuts but also leaves a company that still needs to fund something in the future, and cash burn between now and any wind-down or pivot doesn't stop the day the data reads out. The other asset that would still exist is SLS009 (tambiciclib), the company's CDK9 inhibitor, which has shown positive Phase 2 signals in relapsed/refractory AML and is now dosing a newly diagnosed frontline AML study. It's real, but it's early — Phase 2, unpartnered, competing in a crowded CDK9/cell-cycle inhibitor field — and the kind of risk-adjusted value the market typically assigns to an asset at that stage is modest relative to a name currently carrying a multi-billion-dollar enterprise value. Putting a generous, fully risk-adjusted value on SLS009 optionality on top of the cash gets you to something like \$0.80–\$1.10 per share as a reasonable floor, which lines up closely with a straightforward \$1 ceiling estimate on the combined cash-plus-SLS009 downside case. Against a \$14.98 starting price, that's roughly 93–95% downside in the failure scenario, which, at a 5% probability of success, should be doing most of the work in anyone's expected-value calculation on this stock today.

None of this is a certainty. It's a case built on base rates, mechanism, and pattern-matching against a lot of prior peptide-vaccine history — exactly the kind of case that can be wrong in an individual instance even when the underlying reasoning is sound. But I think it's a stronger, better-evidenced case than “the trial is taking a long time, so it must be working,” and I think the current valuation leaves very little room for being wrong in the market's favor.

You can find my model in my GitHub: <https://github.com/nekkess>

Sources

- Synapse/Patsnap, “How many FDA approved Synthetic peptide vaccine are there?” (2025).
<https://synapse.patsnap.com/article/how-many-fda-approved-synthetic-peptide-vaccine-are-there>
- ACS Biochemistry, “From Neoantigens to Nanocarriers” (2026).
<https://pubs.acs.org/doi/10.1021/acs.biochem.5c00720>
- PMC, “Advancements and Challenges in Peptide-Based Cancer Vaccination”.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC11359700/>
- Frontiers in Immunology, “Peptide-based vaccine for cancer therapies” (2023).
<https://www.frontiersin.org/journals/immunology/articles/10.3389/fimmu.2023.1210044/full>
- Hindawi / Journal of Oncology, “Development of Peptide-Based Vaccines for Cancer” (2022).
<https://www.hindawi.com/journals/jo/2022/9749363/>
- PMC, “Strategies for developing and optimizing cancer vaccines”.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC6518434/>
- PMC, “Trial watch: Peptide-based vaccines in anticancer therapy”.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC6279318/>
- ScienceDirect Topics, Rindopepimut / ACT IV overview. <https://www.sciencedirect.com/topics/medicine-and-dentistry/rindopepimut>
- Nature Reviews Neurology, news piece on the ACT IV failure.
<https://www.nature.com/articles/nrneurol.2016.38>
- Journal of Clinical Investigation (2004), peptide-MHC instability study.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC503773/>
- Maslak et al., Blood Advances (2018), Phase 2 trial of GPS in AML, CR1 (NCT01266083).
<https://ashpublications.org/bloodadvances/article/2/3/224/15779/Phase-2-trial-of-a-multivalent-WT1-peptide-vaccine>
- PMC mirror of the Blood Advances paper above. <https://pmc.ncbi.nlm.nih.gov/articles/PMC5812332/>
- ClinicalTrials.gov, NCT01266083 study record. <https://clinicaltrials.gov/study/NCT01266083>
- CURE Today, “Galinpepimut-S Immunotherapy Has Preliminary Efficacy in AML”.
<https://www.curetoday.com/view/galinpepimut-s-immunotherapy-has-preliminary-efficacy-in-acute-myeloid-leukemia>
- OncLive, “Maintenance Galinpepimut-S Appears Effective and Safe in AML in Second CR”.
<https://www.onclive.com/view/maintenance-galinpepimut-s-appears-effective-and-safe-in-aml-in-second-cr>
- OncLive, “Galinpepimut-S Earns FDA Rare Pediatric Disease Designation for AML”.
<https://www.onclive.com/view/galinpepimut-s-earns-fda-rare-pediatric-disease-designation-for-aml>
- Targeted Oncology, “Phase 3 REGAL Trial of Galinpepimut-S in AML Advances Toward Completion”.
<https://www.targetedonc.com/view/phase-3-regal-trial-of-galinpepimut-s-in-aml-advances-toward-completion>
- CancerNetwork, “Galinpepimut-S Completes Phase 3 REGAL Interim Analysis in AML”.
<https://www.cancernetwork.com/view/galinpepimut-s-completes-phase-3-regal-interim-analysis-in-aml>
- SELLAS press release (Dec 29, 2025) via Yahoo Finance, REGAL 72-event update.
<https://finance.yahoo.com/news/sellas-life-sciences-provides-pivotal-133500622.html>
- StockTitan, republication of the Dec 29, 2025 SELLAS release. <https://www.stocktitan.net/news/SLS/sellas-life-sciences-provides-update-on-pivotal-phase-3-regal-trial-42cvdnc17bmu.html>
- Stocktwits, coverage of Martin Shkreli’s public bearish REGAL comments (March 2026).
<https://stocktwits.com/news-articles/markets/equity/sls-stock-rallies-ahead-of-aml-trial-readout-martin-shkreli-bearish-predicts-failure/cZdoO6oRleg>

- ESMO Open, “Accrual and statistical power failure in published adjuvant phase III oncology trials: 2013–2023”.
[https://www.esmoopen.com/article/S2059-7029\(24\)01372-3/fulltext](https://www.esmoopen.com/article/S2059-7029(24)01372-3/fulltext)
- ClinicalTrials.gov, S0931 (everolimus, kidney cancer) protocol/SAP. https://cdn.clinicaltrials.gov/large-docs/49/NCT01120249/Prot_SAP_ICF_000.pdf
- ClinicalTrials.gov, head/neck RT-cetuximab trial protocol/SAP. https://cdn.clinicaltrials.gov/large-docs/07/NCT00956007/Prot_SAP_000.pdf
- ClinicalTrials.gov, pancreatic cancer adjuvant trial protocol/SAP. https://cdn.clinicaltrials.gov/large-docs/49/NCT01013649/Prot_SAP_000.pdf
- TCTMD, “Event Rates Often Overestimated in Cardiovascular Trials”. <https://www.tctmd.com/news/event-rates-often-overestimated-cardiovascular-trials>
- PharmaTimes, GenVec TNFerade pancreatic cancer trial failure.
https://pharmatimes.com/news/genvecs_tnferade_fails_in_pancreatic_cancer_trial_981540/
- Fierce Biotech, Novartis NOV-002 Phase 3 discontinued (NSCLC).
<https://www.fiercebiotech.com/biotech/novartis-therapeutics-inc-nvlt-discontinues-current-development-program-for-nov-002-nsclc>
- MDedge, canfosfamide (Telik) disappoints in ovarian cancer trials.
<https://www.mdedge.com/obgyn/article/49572/gynecology/canfosfamide-disappoints-ovarian-cancer-trials>
- BioWorld, FDA clinical hold on Telik's Telcyta after failed Phase 3s.
<https://www.bioworld.com/articles/452838>
- PubMed, Vergote et al., Phase 3 canfosfamide vs. PLD/topotecan, ovarian cancer.
<https://pubmed.ncbi.nlm.nih.gov/19515553/>
- TargetedOnc, pixantrone Phase III B-cell NHL trial and FDA ODAC rejection.
<https://www.targetedonc.com/view/pixantrone-fails-to-prolong-pfs-in-phase-iii-bcell-nhl-trial>
- ClinicalTrials.gov, ALSYMPCA (radium-223, prostate cancer) study record.
<https://clinicaltrials.gov/study/NCT00699751>
- PMC, GPS + nivolumab Phase 1, WT1-expressing ovarian cancer.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC10001251/>
- JTO Clinical and Research Reports, GPS + nivolumab Phase 1, diffuse pleural mesothelioma.
[https://www.jtocrr.org/article/S2666-3643\(24\)00126-7/fulltext](https://www.jtocrr.org/article/S2666-3643(24)00126-7/fulltext)
- Fierce Biotech, “Anavex stock sinks 30% as lead asset fails to improve Rett syndrome in children” (2024).
<https://www.fiercebiotech.com/biotech/anavex-stock-sinks-after-rare-disease-trial-fail>
- PMC, letter raising concerns about Anavex's blarcamesine Alzheimer's trial reporting.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC12183999/>
- BMF Reports, “Anavex Life Sciences: How Anavex Built a Story Without a Drug”.
<https://bmfreports.com/articles/anavex>
- PubMed, Macfarlane et al., blarcamesine ANAVEX2-73-AD-004 Phase IIb/III trial results.
<https://pubmed.ncbi.nlm.nih.gov/39800452/>

Note: valuation and TAM figures reference the author's companion Excel model (share price, share count, cash, EV, and peak-sales build-out), current as of July 5, 2026. Stock borrow rate is a market data point subject to change and should be re-verified at the time of reading.